

Opulence Mainframe Deep Research Agent Architecture

1. Simple System Overview

The Opulence system takes legacy mainframe code from a **private wealth bank's security transaction processing system** and makes it understandable using modern AI technology:

- **Input:** COBOL programs, JCL job scripts, PROC procedures, DB2 database definitions, and sample transaction data files from the bank's security trading platform
- **Processing:** Parses and loads them into structured format using a code parser and data loader
- **AI Analysis:** Uses a GPU-hosted CodeLLaMA model (exposed via HTTP API) to analyze and summarize complex business logic
- **Orchestration:** A Coordinator Agent manages the workflow across various specialized research agents
- **Output:** Provides lineage maps showing how customer data flows, business logic summaries explaining trading rules, comprehensive documentation, and an interactive chat interface for asking questions

Example Scenario: Understanding how a customer's security purchase order flows through 50+ COBOL programs, what validation rules apply, and how it updates the portfolio database.

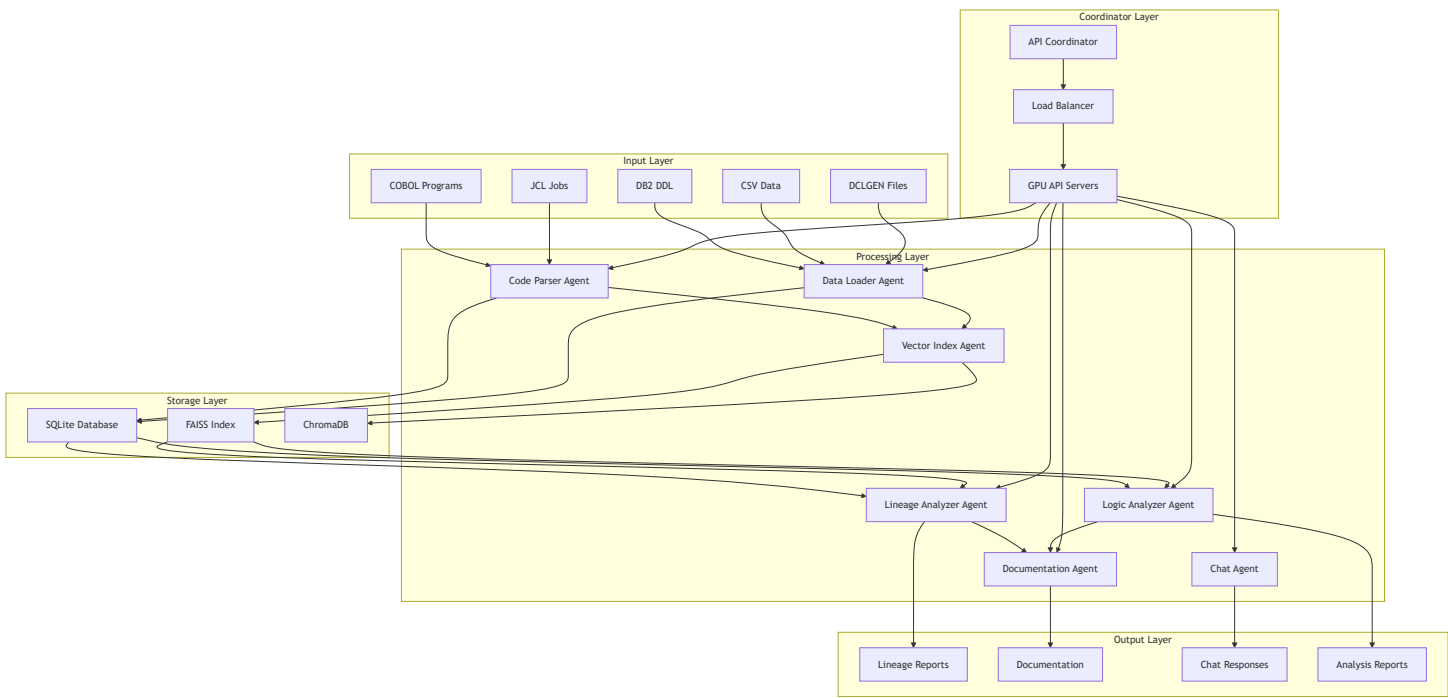
2. Core Components

Component	Function	Value
Code Parser	Converts COBOL/JCL into structured AST	Enables structured understanding of 40-year-old logic
Data Loader	Loads DB2 tables and sample transaction files	Adds real-world context from actual customer trades
Vector Index Agent	Embeds and indexes all elements in FAISS	Powers fast semantic search: "find all margin calculation logic"

Component	Function	Value
Lineage Agent	Tracks fields across jobs and programs	Critical for compliance: trace customer ID through entire system
Logic Analyzer Agent	Extracts business logic and conditional rules	Automates discovery of trading rules and validation logic
Documentation Agent	Summarizes components and logic	Generates readable docs explaining arcane settlement processes
Chat Agent	Interfaces with user to answer questions	"How does stop-loss order processing work?" gets instant answers
Coordinator Agent	Orchestrates flow and agent sequencing	Ensures systematic analysis of interconnected trading systems
GPU LLM API	CodeLLaMA exposed via API for summarization	Core intelligence for understanding legacy financial code

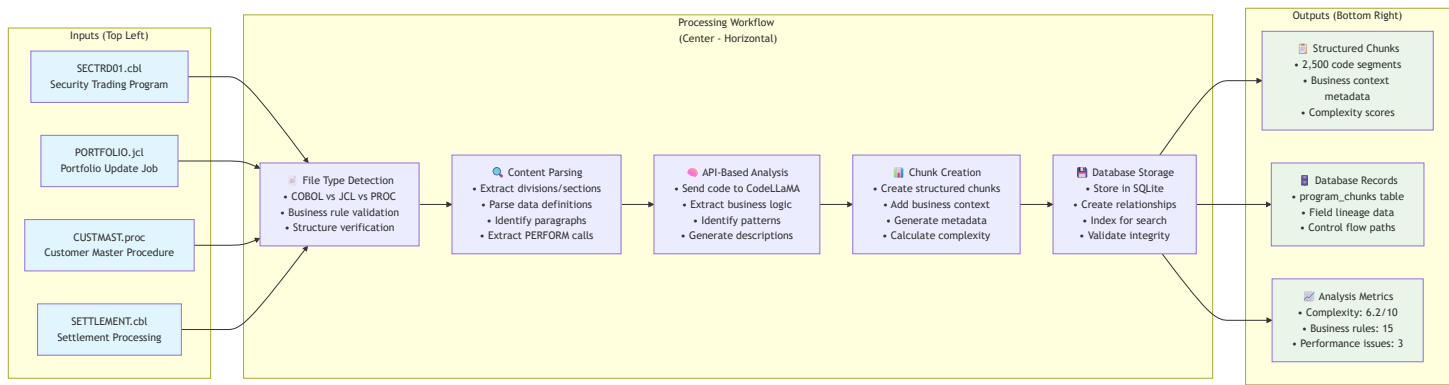
3. System Flow and Individual Agent Workflows

Overall System Architecture Flow

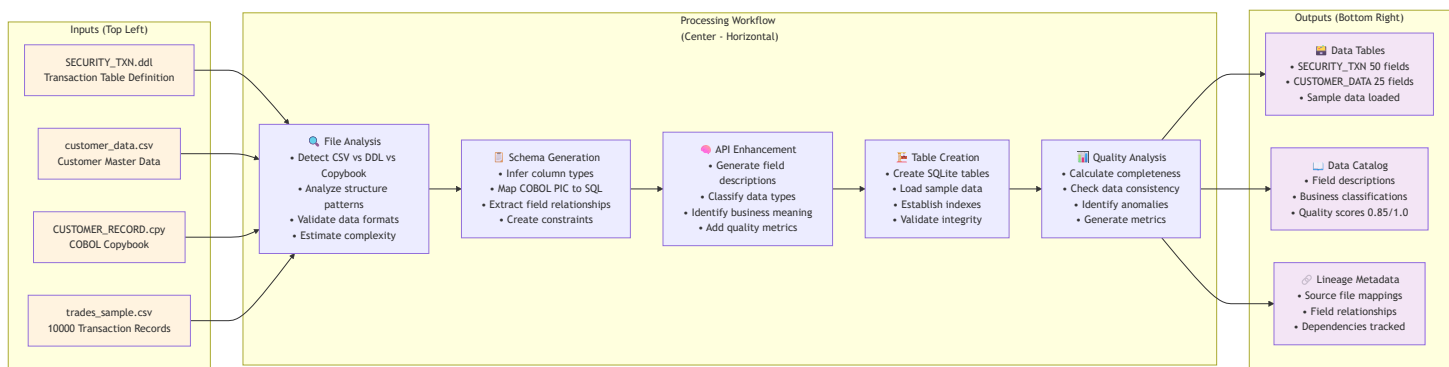


4. Individual Agent Workflows

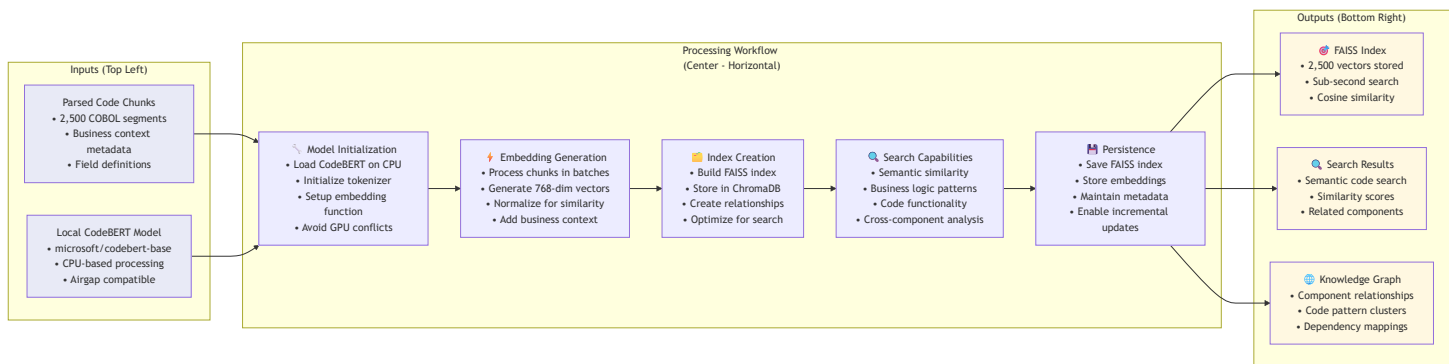
4.1 Code Parser Agent Flow



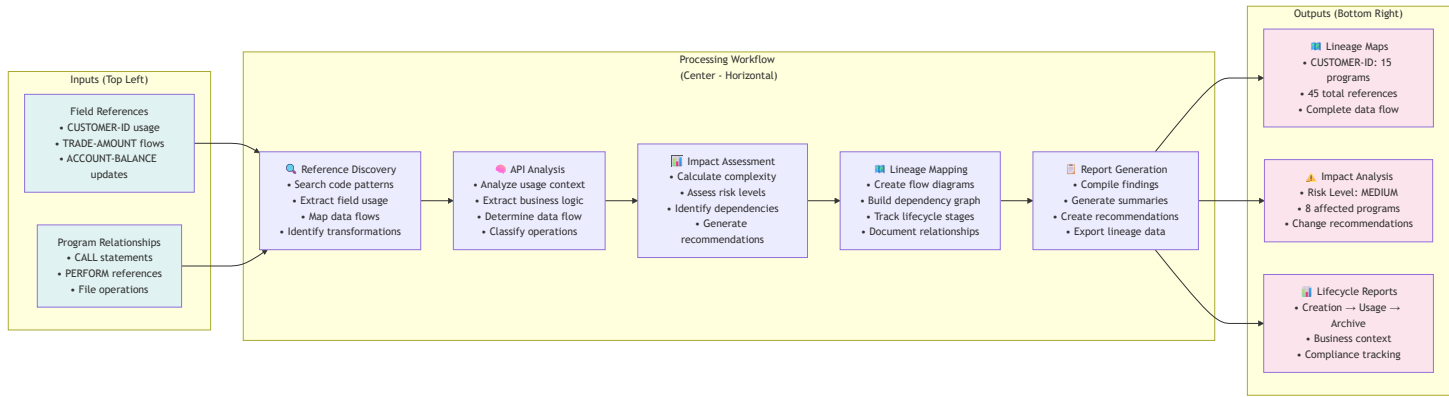
4.2 Data Loader Agent Flow



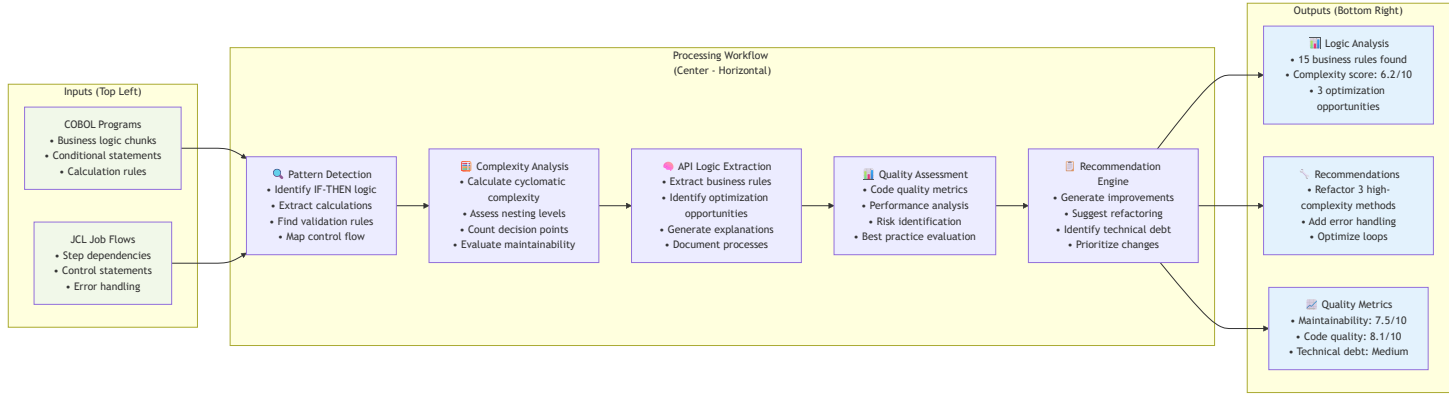
4.3 Vector Index Agent Flow



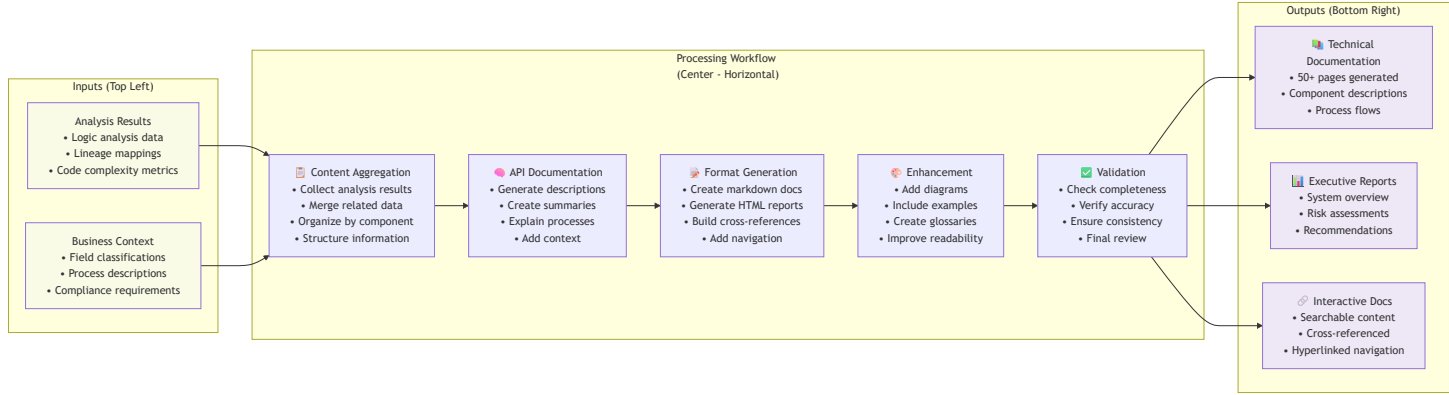
4.4 Lineage Analyzer Agent Flow



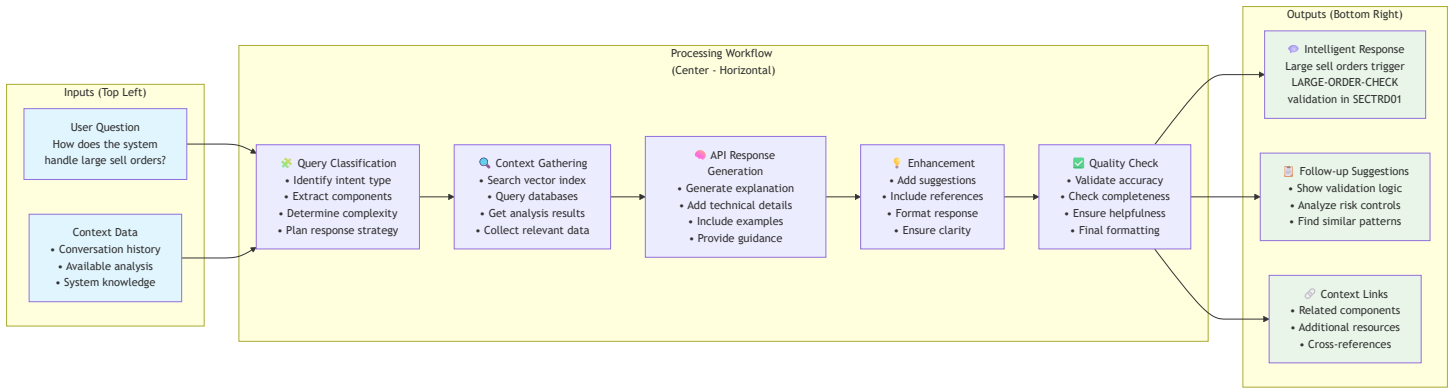
4.5 Logic Analyzer Agent Flow



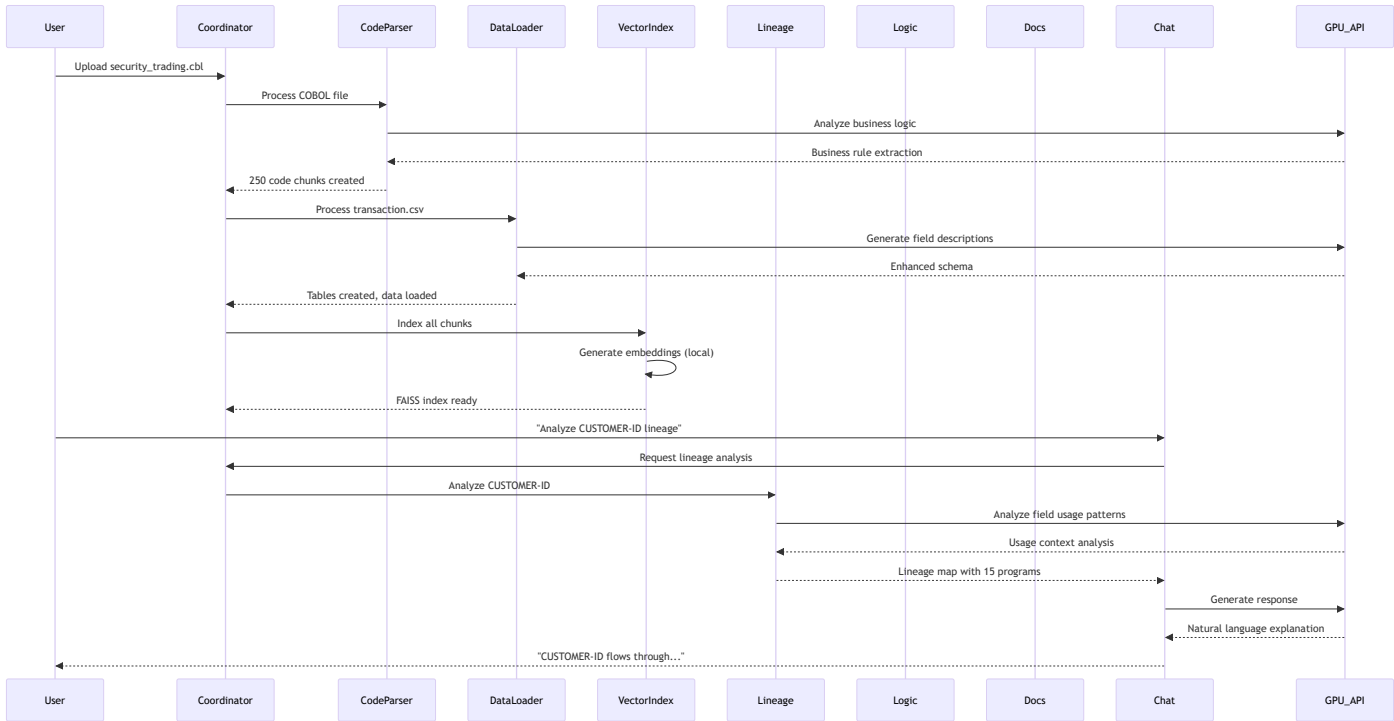
4.6 Documentation Agent Flow



4.7 Chat Agent Flow



5. Agent Coordination Flow



6. Output Artifacts

The Opulence system produces these deliverables for the bank's security trading system:

- ✔ Field-level data lineage reports

- "CUSTOMER-ID flows from CUSTMAST → SECTRD01 → PORTFOLIO-UPDATE → TRADE-HISTORY"
- Compliance-ready audit trails

✓ **Extracted business logic summaries**

- "Stop-loss orders: IF CURRENT-PRICE < (STOP-PRICE * 0.95) THEN EXECUTE-SELL"
- Trading rule documentation in plain English

✓ **Annotated markdown documentation of code modules**

- Complete explanation of settlement processing
- Cross-references between related programs

✓ **Interactive chat interface for querying understanding**

- "What happens when a trade fails settlement?"
- "Show me all programs that update customer portfolios"

7. Sample Data Context: Private Wealth Bank Security Transactions

Input Files for Analysis:

COBOL Programs:

- SECTRD01.cb1 - Main security trading program (2,500 lines)
- VALIDATE.cb1 - Order validation logic (800 lines)
- SETTLE.cb1 - Settlement processing (1,200 lines)
- PORTFOLIO.cb1 - Portfolio update logic (900 lines)

JCL Jobs:

- DAILYTRD.jc1 - Daily trade processing batch job
- SETTLEMENT.jc1 - End-of-day settlement job
- RECON.jc1 - Trade reconciliation job

DB2 Tables:

```
-- SECURITY_TRANSACTION table
CREATE TABLE SECURITY_TXN (
    CUST_ID          CHAR(10),
    TRADE_ID         CHAR(15),
    SECURITY_CODE     CHAR(8),
    TRADE_TYPE       CHAR(4),    -- BUY/SELL
    QUANTITY         DECIMAL(15,2),
    PRICE            DECIMAL(15,4),
    TRADE_DATE       DATE,
    SETTLE_DATE      DATE,
    STATUS           CHAR(3)     -- PEN/SET/FAI
);
```

Sample Transaction Data:

```
CUST_ID,TRADE_ID,SECURITY_CODE,TRADE_TYPE,QUANTITY,PRICE,TRADE_DATE,STATUS
PWB0001234,TRD20241201001,AAPL,BUY,100,150.25,2024-12-01,PEN
PWB0001234,TRD20241201002,TSLA,SELL,50,245.80,2024-12-01,SET
PWB0001567,TRD20241201003,MSFT,BUY,200,380.15,2024-12-01,FAI
```

8. Individual Agent Explanations

Vector Index Agent

Purpose: Creates searchable embeddings of all code segments and business logic.

Bank Example: When analyzing the security trading system, this agent:

- Embeds all COBOL paragraphs dealing with order validation
- Creates vectors for trading rule conditions
- Enables semantic search like "find all margin calculation logic"

API Integration: Makes HTTP calls to CodeLLaMA to generate embeddings and understand code semantics.

Lineage Agent

Purpose: Tracks how data fields flow through the entire system.

Bank Example: For a customer security purchase:

1. **CUSTOMER-ID** enters via online trading platform
2. Flows through `VALIDATE.cb1` for KYC checks
3. Processed in `SECTRD01.cb1` for order execution
4. Updates `PORTFOLIO.cb1` for position management
5. Records in `TRADE-HISTORY` table for audit

Critical for Compliance: Regulators require complete audit trails showing how customer data is processed.

Logic Analyzer Agent

Purpose: Extracts and explains complex business rules embedded in COBOL logic.

Bank Example: Discovers trading rules like:

```
IF TRADE-AMOUNT > DAILY-LIMIT
  AND CUSTOMER-TIER NOT = 'PLATINUM'
  THEN MOVE 'HOLD' TO TRADE-STATUS
  PERFORM MANUAL-APPROVAL-PROCESS
```

Translates to: "Trades over daily limit require manual approval unless customer is Platinum tier."

Documentation Agent

Purpose: Creates human-readable documentation explaining system functionality.

Bank Example: Generates documentation like:

- "Settlement Process Overview: How T+2 settlement works"
- "Stop-Loss Order Processing: Automated selling when price thresholds are breached"
- "Customer Portfolio Updates: Real-time vs. batch processing logic"

Chat Agent

Purpose: Provides conversational interface for querying system knowledge.

Bank Example Queries:

- "How does the system handle partial fills on large orders?"
- "What validation checks are performed before executing a trade?"

- "Show me the settlement process for international securities"

Response Example: "When a large order cannot be filled completely, the PARTIAL-FILL-HANDLER in SECTRD01 splits it into smaller chunks and processes them separately, updating the customer's available cash after each partial execution..."

9. Coordination Flow: Processing a Security Transaction

Real-World Scenario: Customer Places \$500K Apple Stock Purchase

1. File Processing Phase:

- Code Parser analyzes SECTRD01.cb1 and extracts order processing logic
- Data Loader imports recent Apple trading data and customer portfolio info
- System identifies all programs involved in large order processing

2. Analysis Phase:

- **Vector Index Agent:** Finds all code segments related to large order handling
- **Lineage Agent:** Maps how customer cash balance flows through the system
- **Logic Analyzer:** Extracts validation rules for large orders (credit checks, position limits)
- **Documentation Agent:** Summarizes the complete order-to-settlement workflow

3. Query Phase:

- Risk manager asks: "What approvals are needed for this trade size?"
- Chat Agent searches indexed knowledge and responds: "Orders over \$250K require senior trader approval per LARGE-ORDER-CHECK paragraph, plus real-time margin calculation..."

4. Compliance Phase:

- Lineage reports show complete audit trail
- Logic summaries document all decision points
- Documentation provides regulatory-compliant process descriptions

This architecture transforms decades-old, undocumented mainframe code into an accessible, searchable knowledge base that supports both operational teams and regulatory compliance requirements.

10. Technical Implementation Notes

API-Based Architecture

The Opulence system uses HTTP APIs to communicate with GPU-hosted CodeLLaMA models, enabling:

- **Scalability:** Multiple model servers can handle concurrent analysis requests
- **Load Balancing:** Requests are distributed across available GPU resources
- **Fault Tolerance:** Circuit breakers and retry logic ensure robust operation
- **Resource Efficiency:** No need for local GPU allocation per agent

Database Design

SQLite database stores:

- **program_chunks:** Parsed code segments with metadata
- **field_lineage:** Data flow tracking for compliance
- **vector_embeddings:** FAISS index references for semantic search
- **processing_stats:** Performance monitoring and audit trails

This architecture enables financial institutions to understand and maintain critical legacy systems while meeting modern regulatory and operational requirements.